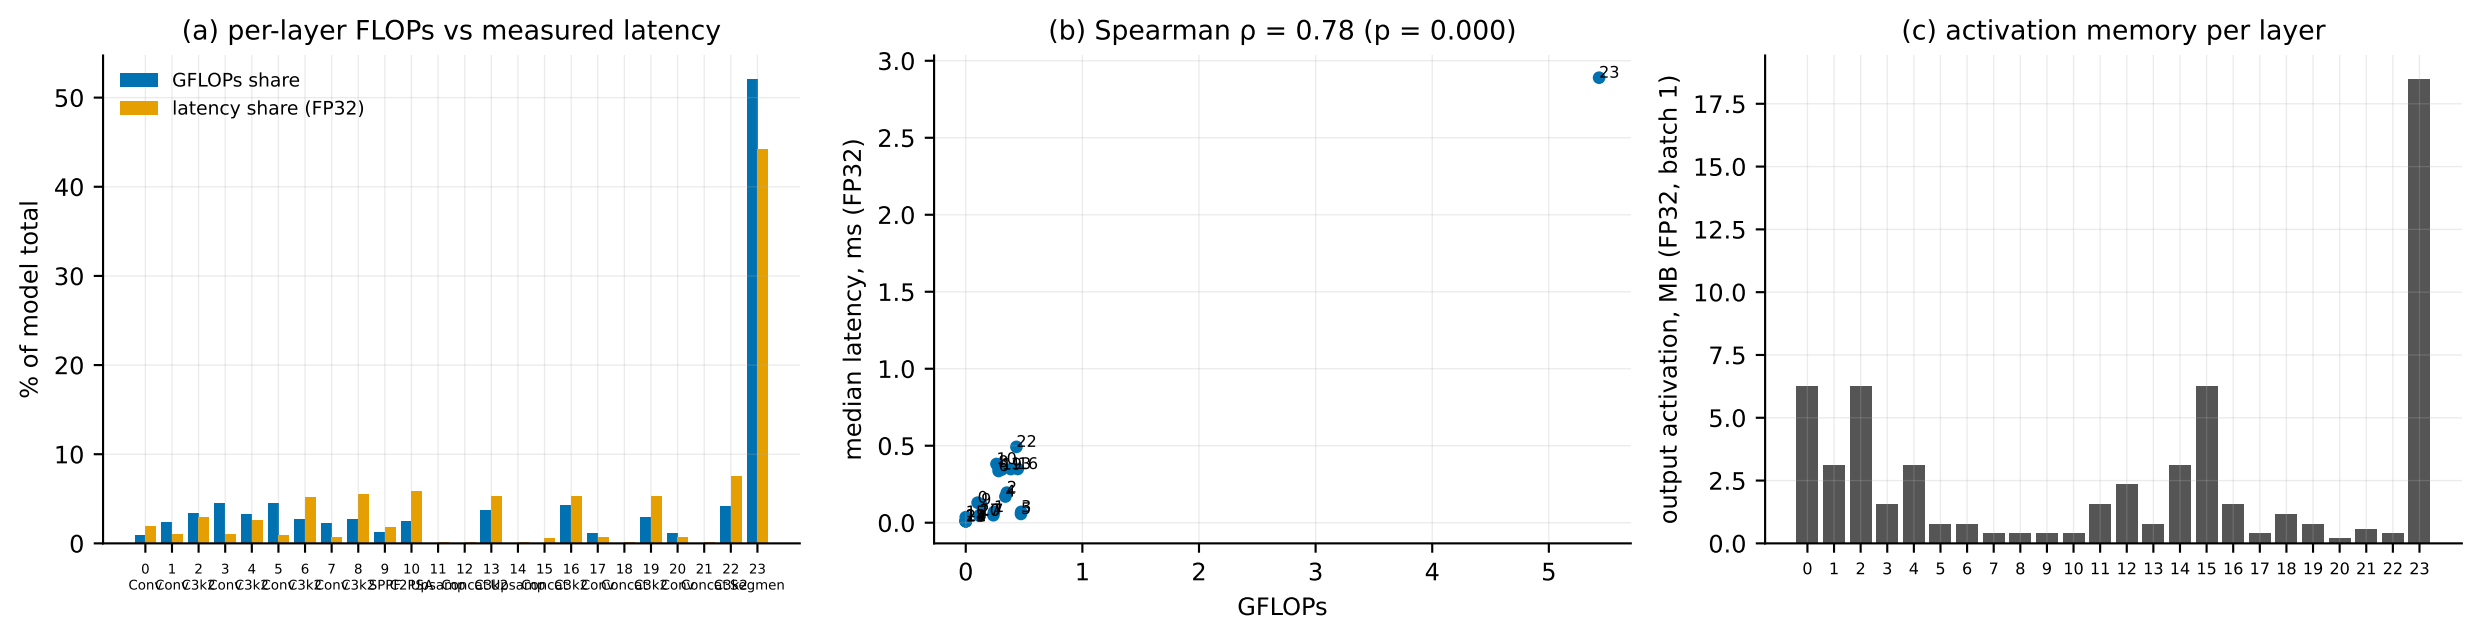
(a) per-layer FLOPs vs measured latency
GFLOPs share
latency share (FP32)
% of model total
(b) Spearman ρ = 0.78 (p = 0.000)
median latency, ms (FP32)
GFLOPs
(c) activation memory per layer
output activation, MB (FP32, batch 1)